

Configuration Manual

MSc Research Project
Data Analytics

Forename Surname
Student ID: x23407719

School of Computing
National College of Ireland

Supervisor: Jorge Basilio

National College of Ireland
MSc Project Submission Sheet
School of Computing

Student Name: Saleem Pasha Shaik

Student ID: 23407719

Programme: MSc Data Analytics

Year: Jan 2025

Module: MSc Research Practicum Part

Lecturer: Jorge Basilio

Submission Due Date: 11/12/2025

Date:

Project Title: Innovative Use of Behavioural Data and Explainable AI for Early Gambling Addiction Detection

Word Count: 1178

Page Count: 11

I hereby certify that the information contained in this (my submission) is information pertaining to research I conducted for this project. All information other than my own contribution will be fully referenced and listed in the relevant bibliography section at the rear of the project.

ALL internet material must be referenced in the bibliography section. Students are required to use the Referencing Standard specified in the report template. To use other author's written or electronic work is illegal (plagiarism) and may result in disciplinary action.

Signature: Saleem Pasha Shaik

Date: 11/12/2025

PLEASE READ THE FOLLOWING INSTRUCTIONS AND CHECKLIST

| | |
|---|--------------------------|
| Attach a completed copy of this sheet to each project (including multiple copies) | <input type="checkbox"/> |
| Attach a Moodle submission receipt of the online project submission, to each project (including multiple copies). | <input type="checkbox"/> |
| You must ensure that you retain a HARD COPY of the project, both for your own reference and in case a project is lost or mislaid. It is not sufficient to keep a copy on computer. | <input type="checkbox"/> |

Assignments that are submitted to the Programme Coordinator Office must be placed into the assignment box located outside the office.

| |
|----------------------------------|
| Office Use Only |
| Signature: |
| Date: |
| Penalty Applied (if applicable): |

Configuration Manual

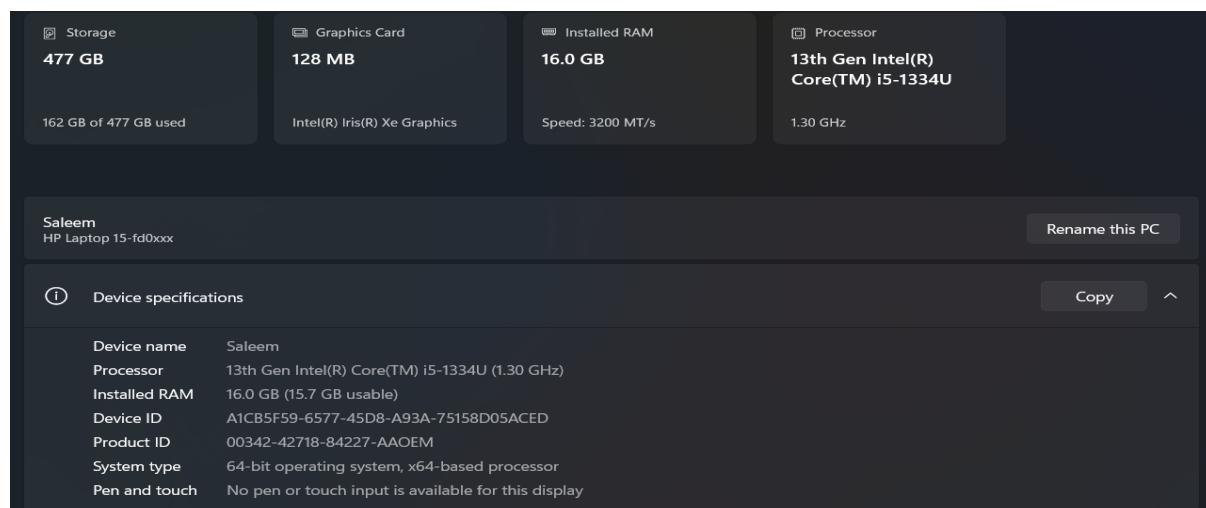
Saleem Pasha Shaik
Student ID: x23407719

1 Introduction

This configuration manual aims to describe the complete system setup used for the implementation and evaluation of the gambling addiction detection framework. The hardware resources, the software environment, the Python packages, and the project structure are defined, allowing other researchers and practitioners to replicate the modelling pipeline correctly. The manual also describes the structure of the project directory, how the dataset is organized, and the flow of execution with respect to every module used in the analysis. Additionally, it provides instructions on how to run the clustering algorithms, train the forecasting models, and explain the AI outputs using LIME, in support of full reproducibility of the results. To maintain accuracy, this manual outlines the exact library versions and dependencies used during experimentation. In this configuration manual, any user will be able to set up the system, execute the complete pipeline, and achieve results that align with those reported in the study.

2 System Configuration

Setting up the environment correctly is crucial for ensuring reproducibility and correct execution of the notebook and model training pipeline. This section details the computing environment, including the hardware and software utilized for running the study and develop. As from the figure, the computational setup consisted of an Intel i5 13th-generation CPU, 16GB RAM, a 64-bit Windows 11 environment, and a 512GB SSD for storage.



| Windows specifications | | Copy | ^ |
|----------------------------------|--|------|---|
| Edition | Windows 11 Home Single Language | | |
| Version | 24H2 | | |
| Installed on | 15-02-2025 | | |
| OS build | 26100.6899 | | |
| Experience | Windows Feature Experience Pack 1000.26100.253.0 | | |
| Microsoft Services Agreement | | | |
| Microsoft Software License Terms | | | |

2.1 Recommended Hardware

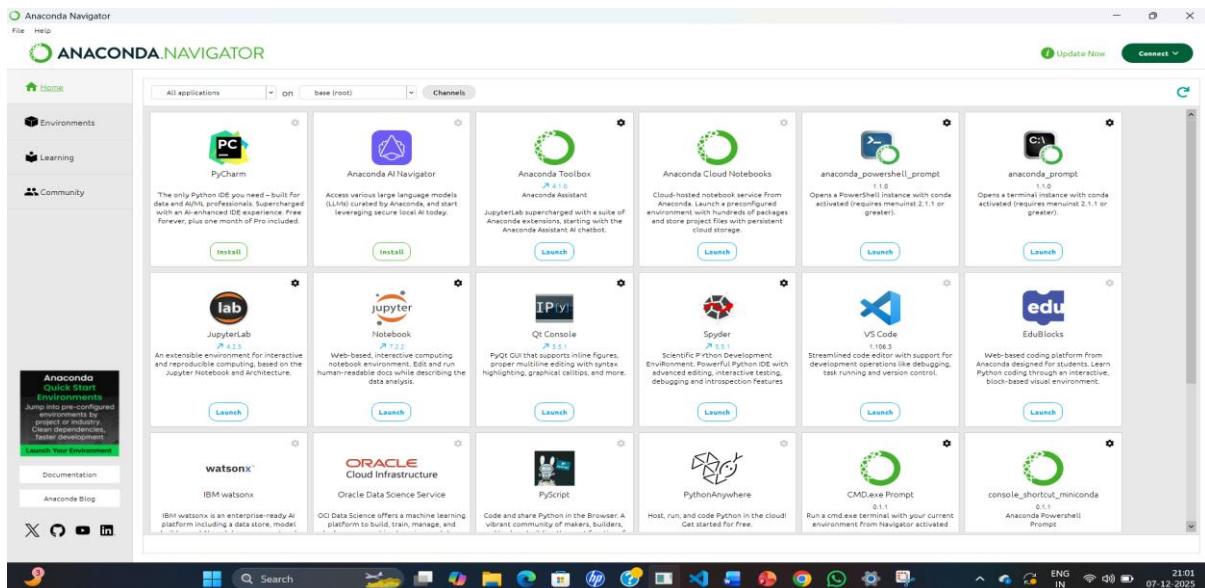
| Component | Recommended |
|-----------|---|
| GPU | NVIDIA RTX 3060 / T4 / A100 (training benefits from CUDA) |
| RAM | Minimum 16 GB, recommended 32 GB |
| CPU | Quad-core or better |
| Storage | 10–20 GB for datasets & models |

Deep learning models, especially attention-based ones, perform significantly faster on a dedicated GPU.

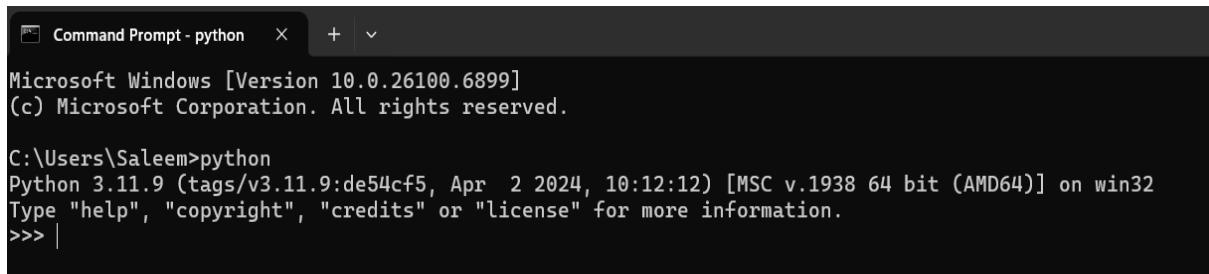
2.2 Software Requirements

Operating System

- Windows 10/11
- Anaconda Navigator
- Jupiter Notebook 7.2.2
- Linux (Ubuntu highly recommended)



2.3 Python Version



```
Command Prompt - python  +  Microsoft Windows [Version 10.0.26100.6899]
(c) Microsoft Corporation. All rights reserved.

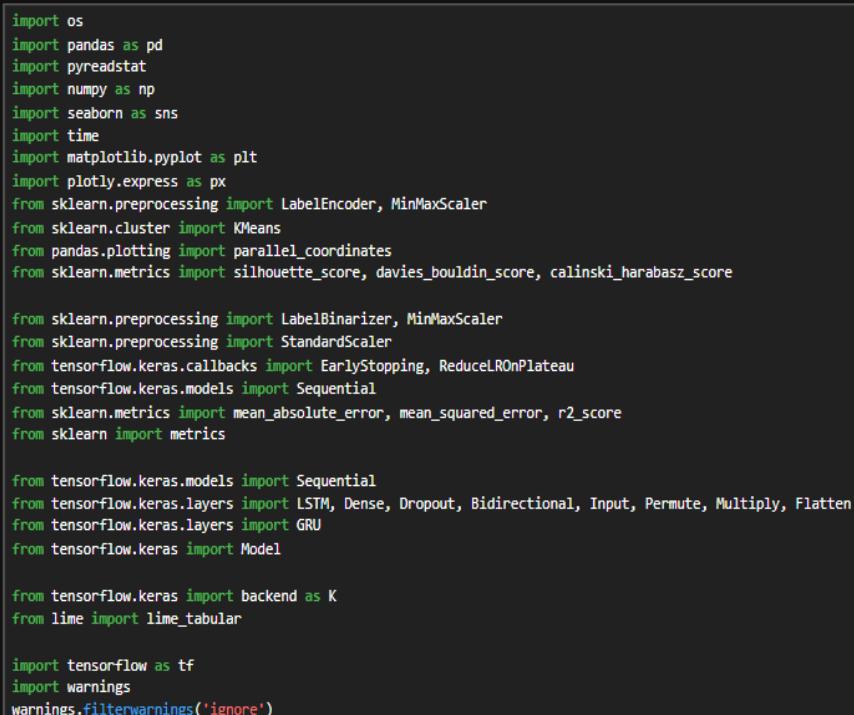
C:\Users\Saleem>python
Python 3.11.9 (tags/v3.11.9:de54cf5, Apr 2 2024, 10:12:12) [MSC v.1938 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license" for more information.
>>> |
```

From above screenshot, we can see the version of python which is used for this project. Python 3.10 or 3.11 recommended

2.1 Required Libraries

The project uses the following major libraries:

| Category | Libraries Used |
|----------------------------------|--|
| Core Scientific Libraries | numpy, pandas, matplotlib, seaborn, scipy, statsmodels |
| Machine Learning / Deep Learning | scikit-learn, tensorflow, keras, keras_cv_attention_models, joblib |
| Explainable AI | lime |
| System Utilities | os, warnings, datetime, time |
| Visualization | matplotlib.pyplot, seaborn |



```
import os
import pandas as pd
import pyreadstat
import numpy as np
import seaborn as sns
import time
import matplotlib.pyplot as plt
import plotly.express as px
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.cluster import KMeans
from pandas.plotting import parallel_coordinates
from sklearn.metrics import silhouette_score, davies_bouldin_score, calinski_harabasz_score

from sklearn.preprocessing import LabelBinarizer, MinMaxScaler
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.models import Sequential
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn import metrics

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout, Bidirectional, Input, Permute, Multiply, Flatten
from tensorflow.keras.layers import GRU
from tensorflow.keras import Model

from tensorflow.keras import backend as K
from lime import lime_tabular

import tensorflow as tf
import warnings
warnings.filterwarnings('ignore')
```

3 Data Sources & Linking

The project uses three datasets: demographics, betting history, and responsible-gambling interventions. The datasets chosen for the project Transparency Project and the title of the data is: Behavioral Characteristics of Internet Gamblers Who Trigger Corporate Responsible Gambling Interventions² (1) . All datasets are merged using “UserID” to create unified player records. An **outer join** ensures no player is excluded

```
# Load SPSS .sav file
df1, meta = pyreadstat.read_sav("C:\\\\Users\\\\Saleem\\\\Downloads\\\\Raw Dataset I.Demographics_Gray_LaPlante_PAB_2012.sav")
df2, meta = pyreadstat.read_sav("C:\\\\Users\\\\Saleem\\\\Downloads\\\\Raw Dataset II.Daily aggregates_Gray_LaPlante_PAB_2012 (1).sav")
df3, meta = pyreadstat.read_sav("C:\\\\Users\\\\Saleem\\\\Downloads\\\\Raw Dataset III.Responsible gambling details_Gray_LaPlante_PAB_2012.sav")
```

```
merged_df = df1.merge(df2, on='UserID', how='outer')
merged_df = merged_df.merge(df3, on='UserID', how='outer')
```

4 Data Cleaning & Feature Engineering

Missing values, inconsistent timestamps, and invalid numeric fields are corrected. Additional features like turnover mean, volatility, intervention counts, and recency are engineered. These enriched features support meaningful clustering.

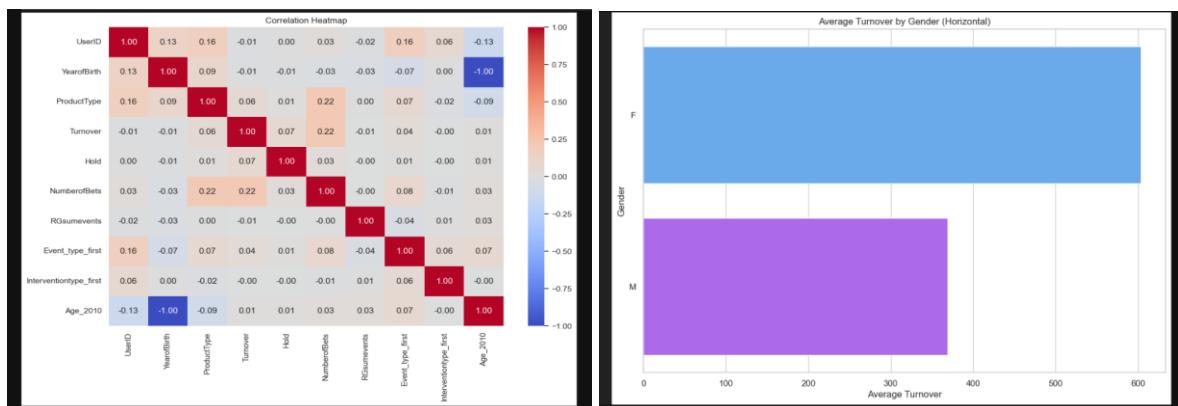
```
# Numeric columns + fill with mean
merged_df.fillna(merged_df.mean(numeric_only=True), inplace=True)

merged_df.dropna(inplace=True)

# FEATURE ENGINEERING
merged_df['Age_2010'] = 2010 - merged_df['YearofBirth']
df = merged_df.copy()
```

5 Data Visualization

Basic visual analysis is performed to understand behavioural patterns. Plots include distributions, correlations, time-series trends, and intervention frequencies. These visuals reveal anomalies and support early insights.



6 Data Preprocessing

Categorical variables are label encoded while continuous behavioural features are scaled. Invalid or incomplete records are removed to ensure consistent clustering inputs. The final dataset contains clean, numerical features.

```
non_date_object_cols = merged_df.select_dtypes(include=['object']).columns.tolist()
date_cols = ['Registration_date', 'First_Deposit_Date', 'Date', 'RGFirst_Date', 'RGLast_date']
cat = [col for col in non_date_object_cols if col not in date_cols]

le = LabelEncoder()
for i in cat:
    merged_df[i] = le.fit_transform(merged_df[i])
print("Encoding complete")
```

7 Clustering Methodology (K-Means)

K-Means ($k=3$) segments players into casual, moderate-risk, and high-risk groups. Inputs include turnover statistics, volatility, interventions, and recency. Each user receives a cluster label for downstream analysis.

```
ssd = []

for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X_scaled)
    ssd.append(kmeans.inertia_)

plt.figure(figsize=(12,8))
plt.plot(range(1, 11), ssd, marker='o')
plt.title('Elbow Method for Optimal K')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('Sum of Squared Distances (SSD)')
plt.grid(True)
plt.show()
```

```
[33]: optimal_k = 3
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
kmeans.fit(X)

cluster_labels = kmeans.labels_
cluster_centers = kmeans.cluster_centers_

plt.figure(figsize=(10, 7))
plt.scatter(X.iloc[:, 0], X.iloc[:, 1], c=cluster_labels, cmap='viridis', s=50, alpha=0.8)
plt.scatter(cluster_centers[:, 0], cluster_centers[:, 1], s=200, marker='X', c='red', label='Cluster Centroids')
plt.title(f'K-Means Clustering with K={optimal_k}')
plt.xlabel('Feature 1 (Scaled)')
plt.ylabel('Feature 2 (Scaled)')
plt.legend()
plt.grid(True)
plt.show()

df['Cluster'] = cluster_labels
merged_df.head()
```

8 Cluster Evaluation Metrics

Silhouette Score, DBI, CH Index, and the Elbow Method are used to validate clustering quality. These metrics help verify separation between risk-based player groups. Higher silhouette and CH, lower DBI indicate good clustering.

```
sil_score = silhouette_score(merged_df, merged_df['Cluster'])
db_score = davies_bouldin_score(merged_df, merged_df['Cluster'])
ch_score = calinski_harabasz_score(merged_df, merged_df['Cluster'])

print("\n[Evaluation Metrics]")
print(f"Silhouette Score: {sil_score:.4f} (higher = better, max=1)")
print(f"Davies-Bouldin Index: {db_score:.4f} (lower = better)")
print(f"Calinski-Harabasz Index: {ch_score:.4f} (higher = better)")
```

9 Filtering High-Risk Players

Only the cluster tagged as **High-Risk** is retained for detailed modelling. A single high-risk user is selected for time-series prediction. Their full behavioural timeline is extracted for analysis.

```
valid_users = dataframe['UserID'].value_counts()[lambda x: x > 102].index

dataframe = dataframe[dataframe['UserID'].isin(valid_users)]

print("Remaining users:", dataframe['UserID'].nunique())
print("Total records:", len(dataframe))
```

10 Time-Series Data Preprocessing

Time gaps, missing turnover values, and irregular timestamps are corrected. Invalid entries are removed to maintain chronological consistency. This ensures clean input for sequence modelling.

```
: object_cols = dataframe.select_dtypes(include=['object']).columns.tolist()

le = LabelEncoder()
for i in object_cols:
    dataframe[i] = le.fit_transform(dataframe[i])

print("[info] Encoding complete!")
dataframe = dataframe.set_index("Date").resample("D").sum().fillna(0)
```

11 Window Rolling (7)

A 7-day sliding window converts the sequence into supervised learning format. Inputs (X) represent past 7 days, while outputs (y) predict the next day. Data is reshaped into (samples, 7, features) format for DL models.

```

def create_rolling_windows(data, target, window_size=7):
    X, y = [], []
    for i in range(len(data) - window_size):
        X.append(data[i:i+window_size])
        y.append(target[i+window_size])
    return np.array(X), np.array(y)

features = dataframe[['CountryName', 'LanguageName', 'Gender', 'Turnover', 'Hold', 'NumberofBets', 'RGsumevents', 'Event_type_first', 'Age_2010']].values

# Example target → risky (1) vs non-risky (0) player label
# Replace with your actual labels
target = (dataframe['Turnover']).values

# Create rolling windows
X, y = create_rolling_windows(features, target, window_size=7) # Generate input-output pairs using a 7-day rolling window

print("X shape:", X.shape)
print("y shape:", y.shape)

```

12 Train-Test

The processed sequences are split into **80% train** and **20% test**. This ensures models learn historical trends and generalize to unseen data. Splitting maintains chronological order.

```

split = int(0.8 * len(X))
X_train, X_test = X[:split], X[split:]
y_train, y_test = y[:split], y[split:]

```

13 Early Stop & LR Scheduler

Early stopping was applied to prevent overfitting by monitoring the model's validation loss during training. If the validation loss did not improve for five consecutive epochs, training was halted automatically. The best-performing model weights were restored to ensure optimal generalization. A learning-rate reduction mechanism was also used to lower the learning rate when progress plateaued, allowing the model to refine learning more effectively.

```

early_stop = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True,
    verbose=1
)

lr_reduce = ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.5,
    patience=2,
    verbose=1
)

```

14 Time-Series Model

Four models are trained: LSTM, GRU, BiLSTM, and BiLSTM with Cross Attention. Early stopping and LR scheduling improve stability and reduce overfitting. The best-performing architecture is selected for final evaluation.

LSTM:

```
model = Sequential([
    LSTM(128, activation='tanh', return_sequences=True, input_shape=(X_train.shape[1], X_train.shape[2])),
    Dropout(0.2),
    LSTM(64, activation='tanh', return_sequences=False),
    Dropout(0.2),
    Dense(32, activation='relu'),
    Dense(1)
])

model.compile(
    optimizer='adam',
    loss='mse',
    metrics=['mae']
)

history = model.fit(
    X_train, y_train,
    validation_data=(X_test, y_test),
    epochs=30,
    batch_size=64,
    verbose=1,
    callbacks=[early_stop, lr_reduce]
)
```

GRU:

```
def build_gru_model(window_size, n_features):
    model = Sequential()
    model.add(GRU(64, return_sequences=True, input_shape=(window_size, n_features)))
    model.add(Dropout(0.7))
    model.add(GRU(64, return_sequences=True))
    model.add(Dropout(0.7))
    model.add(GRU(64, return_sequences=True))
    model.add(Dropout(0.7))
    model.add(GRU(64, return_sequences=False))
    model.add(Dropout(0.7))
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mse', metrics=['mae'])
    return model

model = build_gru_model(X_train.shape[1], X_train.shape[2])
model.summary()
```

BiLSTM:

```
# LSTM
model = Sequential()

# First BiLSTM Layer
model.add(Bidirectional(LSTM(units=128, return_sequences=True, input_shape=(X_train.shape[1], X_train.shape[2]))))
model.add(Dropout(0.3))

# Second BiLSTM Layer
model.add(Bidirectional(LSTM(units=64, return_sequences=True)))
model.add(Dropout(0.7))

# Third BiLSTM Layer
model.add(Bidirectional(LSTM(units=32, return_sequences=False)))
model.add(Dropout(0.7))

# Output Layer
model.add(Dense(1))
```

BiLSTM Cross Attention:

```
INPUT_DIMS = X_train.shape[1]
TIME_STEPS = 1
lstm_units = 128

inputs = Input(shape=(X_train.shape[1], X_train.shape[2]))

# First LSTM Layer
x = Bidirectional(LSTM(lstm_units, return_sequences=True))(inputs)
x = Dropout(0.5)(x)

# Second LSTM Layer
x = Bidirectional(LSTM(64, return_sequences=True))(x)
x = Dropout(0.5)(x)

# Attention layer
attention_mul = cross_attention_layers(x)
attention_mul = Flatten()(attention_mul)

# Output layer
output = Dense(1)(attention_mul)

# Define the model
model = Model(inputs=inputs, outputs=output)
```

15 Evaluation Metrics

Performance is assessed using RMSE, MAE, and MSE. Latency and throughput measure system efficiency during inference. Together, they reflect both predictive accuracy and operational performance.

```
predicted = model.predict(X_test)
print('Mean Squared Error(MSE):', metrics.mean_squared_error(y_test, predicted))
print('Root Mean Squared Error(RMSE):', np.sqrt(metrics.mean_squared_error(y_test, predicted)))
print('Mean Absolute Error(MAE):', metrics.mean_absolute_error(y_test, predicted))

predicted_value = scalers['Turnover'].inverse_transform(predicted)
real_value = scalers['Turnover'].inverse_transform(y_test.reshape(-1, 1))

# LATENCY AND THROUGHPUT
start = time.time()
for i in range(min(1000, len(X_test))):

    _ = model.predict(X_test[i].reshape(1, X_test.shape[1], X_test.shape[2]), verbose=0)
end = time.time()
latency = (end - start) / min(1000, len(X_test))

start = time.time()
_ = model.predict(X_test)
end = time.time()
throughput = len(X_test) / (end - start)

print(f"\nLatency per prediction: {latency * 1000:.4f} ms")
print(f"Throughput: {throughput:.2f} predictions/sec")

actualpred_df = pd.DataFrame({
    "Real": real_value.ravel(),
    "Predicted": predicted_value.ravel()
}, index = datafram.index[-len(real_value):])
actualpred_df.tail()
```

16 Explainability AI

LIME identifies which features and time steps drive the model's predictions. It highlights behavioural triggers contributing to risk increases. This ensures transparency in deep-learning-based addiction detection.

```

explain_index = -1
n_features_to_show = 20
n_samples_in_lime = 5000
html_output = "lime_explanation_gambling.html"

assert 'X_train' in globals() and 'X_test' in globals(), "X_train / X_test not found"
assert X_train.ndim == 3, "X_train should be 3D (samples, window, features)"
window_size = X_train.shape[1]
n_features = X_train.shape[2]

feature_cols = ['CountryName', 'LanguageName', 'Gender', 'Turnover', 'Hold', 'NumberofBets', 'RGsumevents', 'Event_type_first', 'Age_2010'] # base feature
feat_names = []

for t in range(window_size):
    lag = window_size - t
    for f in feature_cols:
        feat_names.append(f"({f})_{t}-{(lag)}")

# --- Flatten data for LIME ---
X_train_flat = X_train.reshape(len(X_train), -1)
X_test_flat = X_test.reshape(len(X_test), -1)

# --- Define predict function for LIME ---
def lime_predict_fn(flat_batch):
    X_batch = flat_batch.reshape((flat_batch.shape[0], window_size, n_features)).astype(np.float32)
    preds = model.predict(X_batch, verbose=0)
    return preds.reshape(-1)

explainer = lime_tabular.LimeTabularExplainer(
    training_data=X_train_flat.astype(np.float32),
    feature_names=feat_names,
    mode='regression',
    discretize_continuous=False
)

flat_instance = X_test_flat[explain_index].astype(np.float32)

```

```

exp = explainer.explain_instance(
    data_row=flat_instance,
    predict_fn=lime_predict_fn,
    num_features=n_features_to_show,
    num_samples=n_samples_in_lime
)

exp_list = exp.as_list()
exp_df = pd.DataFrame(exp_list, columns=['feature', 'weight'])

try:
    inst_window = flat_instance.reshape(window_size, n_features)
    readable = []
    for t in range(window_size):
        lag = window_size - t
        for fi, fname in enumerate(feature_cols):
            flat_index = t * len(feature_cols) + fi
            if flat_index < len(flat_instance):
                scaled_val = flat_instance[flat_index]
                readable.append({'feature': f"({fname})_{t}-{(lag)}", 'scaled_value': scaled_val})
    readable_df = pd.DataFrame(readable)
except Exception as e:
    print(f"\nError creating readable_df: {e}")
    readable_df = None

print("\nTop LIME explanation (feature, weight):")
print(exp_df.head(n_features_to_show).to_string(index=False))

if readable_df is not None:
    print("\nCorresponding scaled instance values (first 30 shown):")
    print(readable_df.head(30).to_string(index=False))

# --- Save interactive HTML ---
exp.save_to_file(html_output)
print(f"\n Saved interactive LIME explanation to: {os.path.abspath(html_output)}")

# --- Visualize in notebook ---
try:
    fig = exp.as_pyplot_figure()
    plt.title("LIME Feature Importance - Gambling Addiction Predictor (BiLSTM)")
    plt.tight_layout()
    plt.show()
except Exception:
    pass

```

Repositories:

GitHub: <https://github.com/saleempasha6969-cyber/Research-Practicum>

1) Data Source: http://www.thetransparencyproject.org/download_index.php

Dataset Links:

1. <http://www.thetransparencyproject.org/datasets/Raw%20Datset%20I.Demographics%20Gray%20LaPlante%20PAB%202012.sav>
2. <http://www.thetransparencyproject.org/datasets/Raw%20Datset%20II.Daily%20Aggregates%20Gray%20LaPlante%20PAB%202012.sav>
3. <http://www.thetransparencyproject.org/datasets/Raw%20Datset%20III.Responsible%20Gambling%20Details%20Gray%20LaPlante%20PAB%202012.sav>

References

Python: <https://www.python.org>

Tensorflow : https://www.tensorflow.org/api_docs/python/tf/keras